

Name - Ranjana kumari

Task 1 - Linear Regression using python scikit learn

In this section we will see how to implement linear regression using python scikit learn to predict the percentage of marks that a student is expected to score based upon the number of hours they studied. This is a simple linear regression task as it

```
In [ ]: # Importing all libraries required in this notebook
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: # Reading data from remote link
url = "http://bit.ly/w-data"
study_data = pd.read_csv(url)
print("Data imported successfully")
#displaying the first 10 data
study_data.head(10)
```

Data imported successfully

Out[2]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25

```
In [3]: # shape of the dataset
study_data.shape
```

Out[3]: (25, 2)

```
In [4]: # datatypes of each attribute
study_data.dtypes
```

Out[4]: Hours float64
Scores int64
dtype: object

```
In [5]: # check for null values
study_data.isnull().any()
```

Out[5]: Hours False
Scores False
dtype: bool

```
In [6]: #summary statistics of the data
study_data.describe()
```


Out[6]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

Data Visualization

Plotting a 2D graph on the data points from our dataset using scatter plot to see the distribution of scores and look for any

```
In [7]: # Plots the distribution of scores
study_data.plot(x='Hours', y='Scores', kind='scatter', grid=True)
plt.title('Hours vs Percentage')
plt.xlabel('Hours Studied')
plt.ylabel('Percentage Score')
plt.show()
```



From the graph above, we can clearly see that there is a positive linear relation between the number of hours studied and percentage of score.

Preparing the data

Dividing the data into inputs (all the attributes except the attribute which needs to be predicted) and output (the attribute which

```
In [8]: # Inputs
X = study_data.iloc[:, :-1].values
# Output
Y = study_data.iloc[:, -1].values
```

Splitting the data

Data is split using train test split from Scikit Learn model_selection module.

```
In [9]: from sklearn.model_selection import train_test_split
```

```
In [10]: # Splits data by 80-20 i.e, 80 % of data for training and 20 % of data for testing
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.20, random_state=0)
```

Training the Algorithm

We are using Linear Regression model algorithm for our problem from Scikit Learn's linear_model module library.

```
In [11]: from sklearn.linear_model import LinearRegression
```


```
In [12]: # Instantiating the model
regressor = LinearRegression()
# Fitting the data for training
regressor.fit(X_train, Y_train)

print('Training done.')
```

Training done.

```
In [13]: # Plotting the regression line
line = regressor.coef_*X+regressor.intercept_

# Plotting for the test data
plt.scatter(X, Y)
plt.plot(X, line);
plt.show()
```



Making Predictions

After training the model, its time to predict the output values using the input test data.

```
In [14]: # Predicting the scores
Y_pred = regressor.predict(X_test)
```

```
In [15]: # Comparing Actual vs Predicted
df = pd.DataFrame({'Actual': Y_test, 'Predicted': Y_pred})
df
```

Out[15]:

	Actual	Predicted
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

```
In [16]: # You can also test with your own data
hours = 9.25
own_pred = regressor.predict([[hours]])
print(f"No of Hours = {hours}")
print(f"Predicted Score = {own_pred[0]}")
```

No of Hours = 9.25
Predicted Score = 93.69173248737538

Evaluating the model

The last step is to evaluate the performance and efficiency of the algorithm. This step becomes useful to compare it with different algorithms working on the same dataset to see which algorithm performs the best. Here, we have chosen the mean

```
In [17]: from sklearn.metrics import mean_absolute_error
```

```
In [18]: print('Mean Absolute Error: ',mean_absolute_error(Y_test, Y_pred))

Mean Absolute Error: 4.183859899002975
```

```
In [ ]:
```